

# Event-driven systems

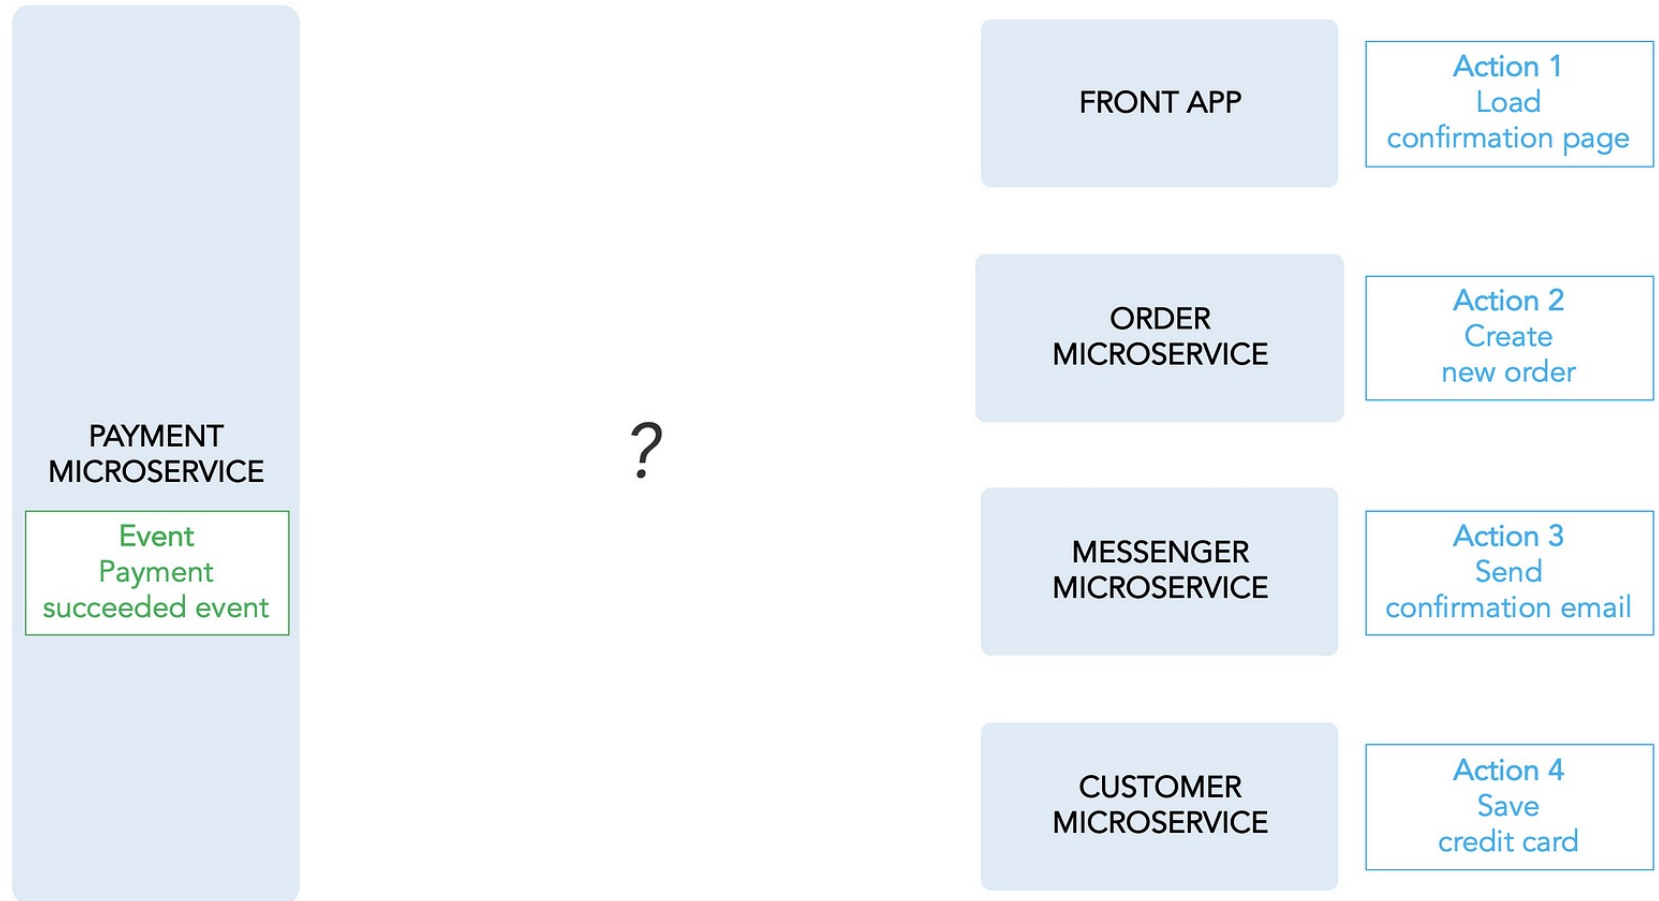
UA.DETI.IES

# Scaling with Microservices

---

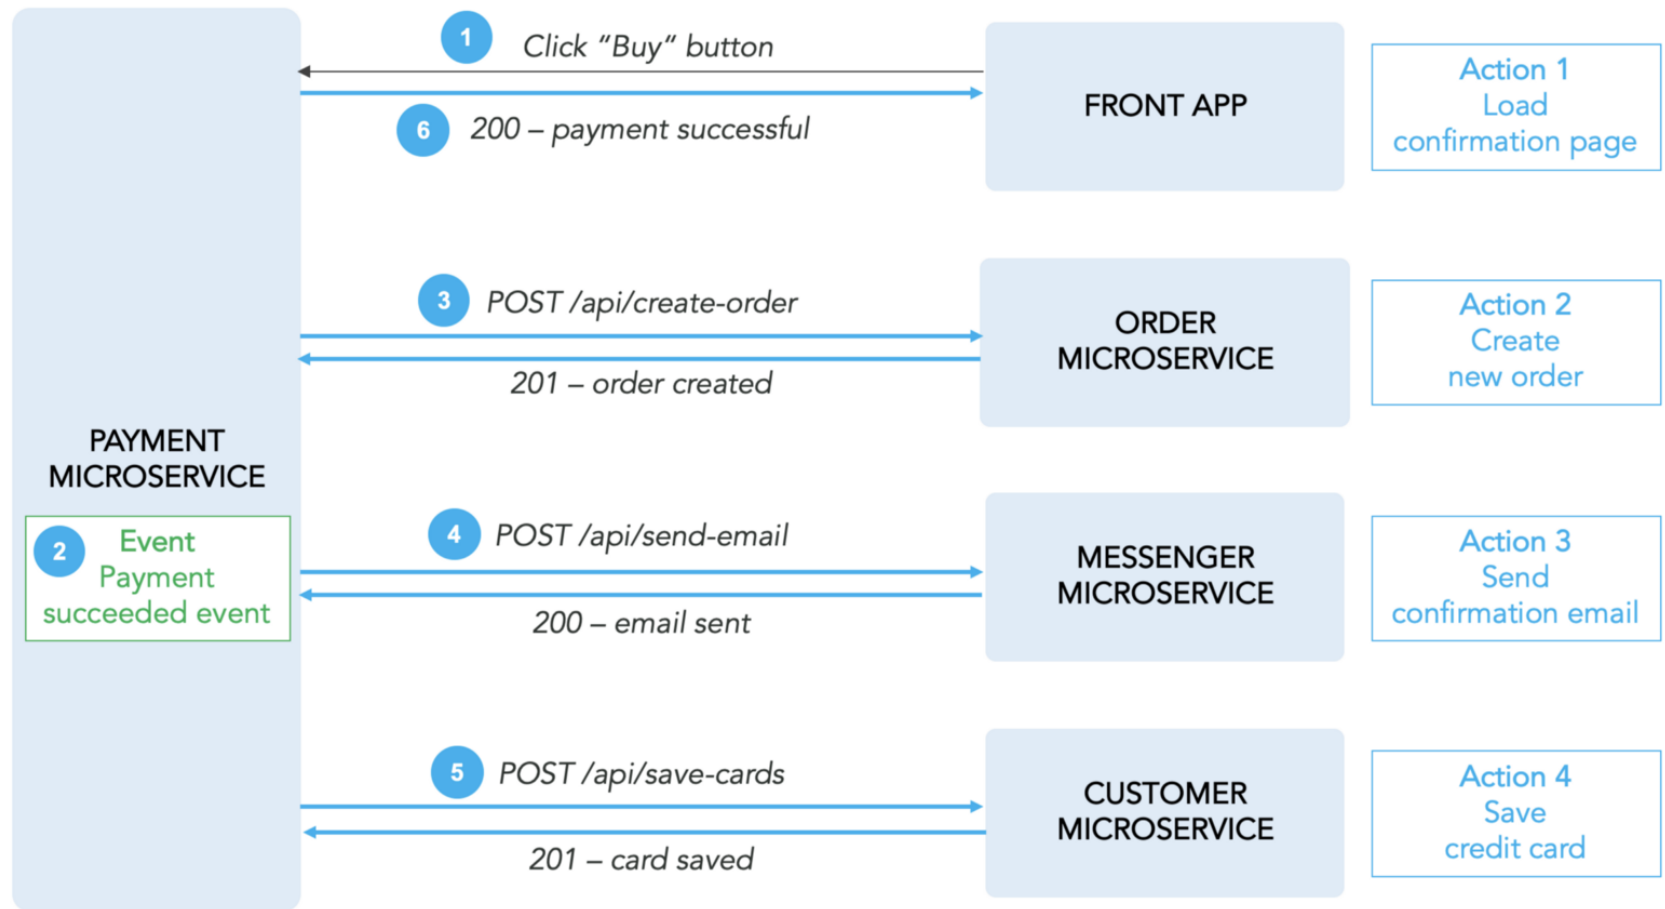
- ❖ **Application development** is currently based on concepts like **microservices** and **serverless** functions
  - We now live in a **cloud-native world**, and these models are a natural fit for distributed cloud-based environments
  - But simply building and deploying services and functions is not enough
- ❖ On its own, a single microservice does not accomplish much
  - **We also need a way to wire up those components**
    - i.e., to connect them so they can exchange data, forming a true application
  - In SE, this is one of the most important architectural decisions to make

# How to notify every Microservice?

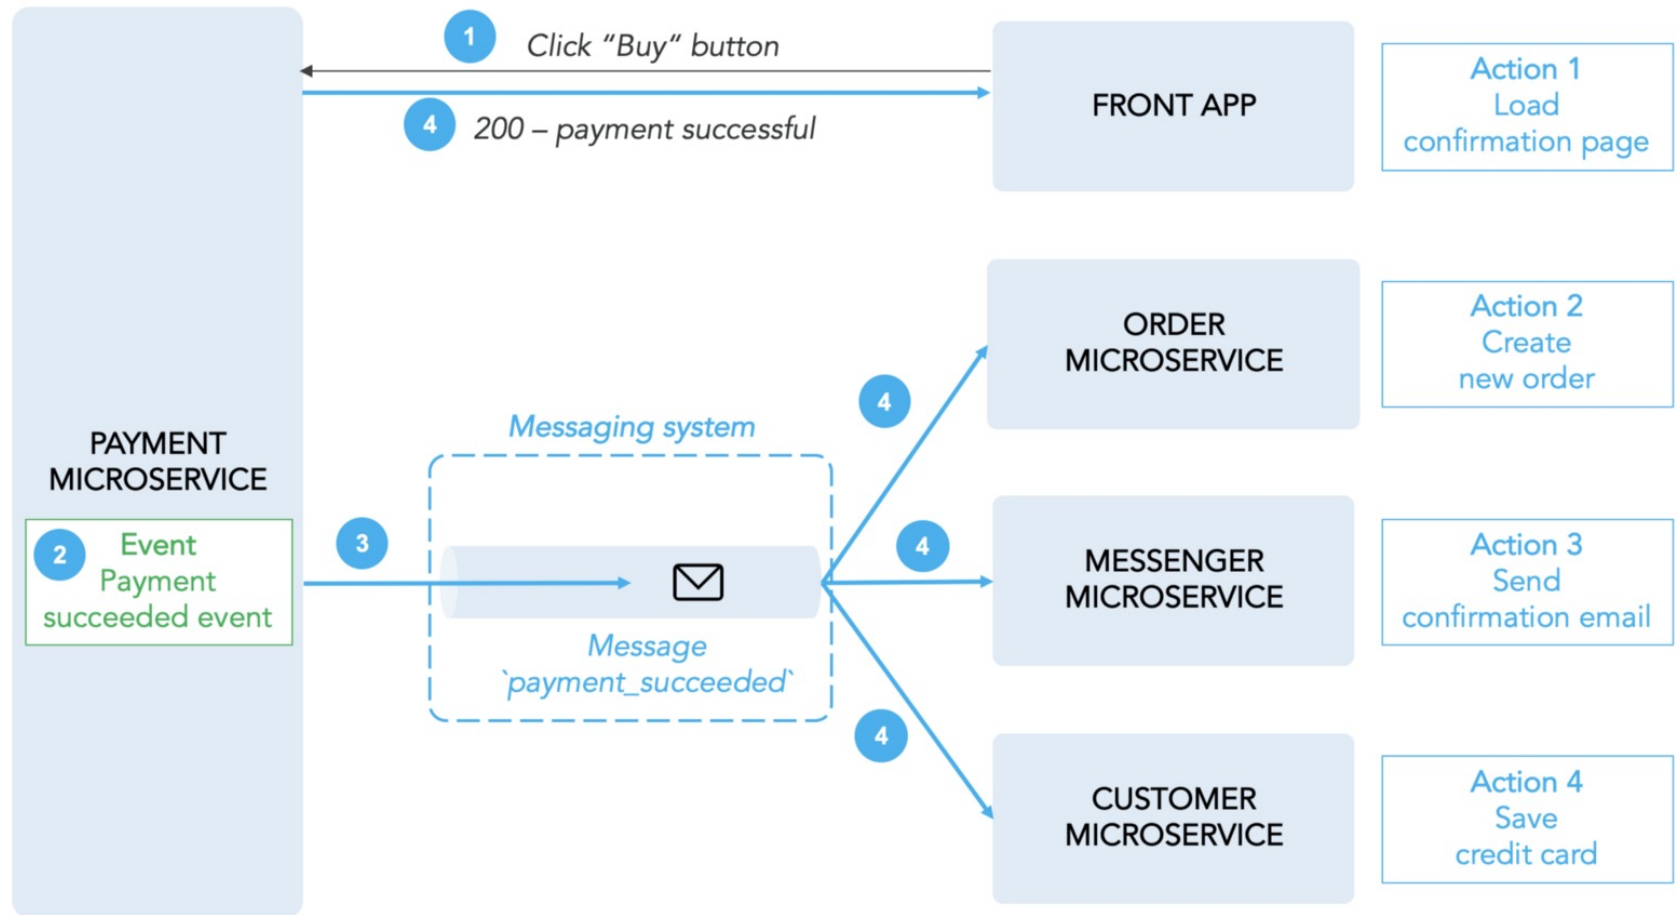


<https://blog.theodo.com/2019/08/event-driven-architectures-rabbitmq/>

# Request-driven architecture



# Event-driven architecture



# Event-driven architecture

---

- ❖ For controlling these services, various mechanisms were developed over the years, such as **message queues** and **enterprise service buses** (ESBs).
  - E.g. **RabbitMQ**, **WSO2**.
- ❖ More recent offerings, the concept of **streaming data** have also emerged.
  - E.g. Apache **Kafka**
- ❖ This latter category is growing
  - Because streaming data is seen as a useful tool for implementing **event-driven architecture**
  - a software design pattern in which application data is modeled as streams of events, rather than as operations on static records.

# What is an event?

---

- ❖ **Events** are things that happen, within a software system or, more broadly, during the operation of a business or other human process.
  - e.g., a sensor reports a temperature change, a user clicks their mouse, a customer deposits a check into a bank account.
- ❖ The concept of events in software systems closely aligns with how most of us think about our day-to-day lives.
  - Organizing around events makes it easier to develop business logic that accurately models real-world processes.
  - It helps reducing the number of one-to-one connections within a distributed system increasing the value of the microservices.
- ❖ An **event-driven architecture** allows generating, storing, accessing and reacting to these events.

# Events ...





# Events vs. queries and commands

---

- ❖ **Queries** are a request to look something up
    - Unlike events or commands, queries are free of side effects; they leave the state of the system unchanged.
  - ❖ **Commands** are actions
    - Requests for some operation to be performed and will change the state of the system.
    - Synchronous and typically indicate completion.
- 

	Behavior/state change	Includes a response
<b>Command</b>	Requested to happen	Maybe
<b>Event</b>	Just happened	Never
<b>Query</b>	None	Always

# Event-driven patterns – notification

---

- ❖ A service sends **events to notify** other systems of a change in its domain
  - For example, a user account service might send a notification event when a new login is created
- ❖ What other systems choose to do with that information is largely up to them
  - The service that issued the notification just carries on with its business
- ❖ Notification events **usually do not carry much data**
  - Resulting in a loosely coupled system with minimal network traffic spent on messaging

# Event-driven patterns – state transfer

---

- ❖ A step up from simple notification, in this model the recipient of an **event also receives the data** it needs to perform further work
  - E.g., the user account service might issue an event that includes a data packet containing the new user's login ID, full name, hashed password, and other pertinent details.
- ❖ This model can be appealing to developers familiar with RESTful interfaces.
  - But, depending on the complexity of the system, it can lead to a lot of data traffic on the network and data duplication in storage

# Event-driven patterns – sourcing

---

- ❖ The goal of **event-sourcing** is to represent every change of state in a system as an event, each recorded in chronological order
- ❖ In so doing, the **event stream itself becomes the principal source of truth for the system**
  - E.g., it should be possible to “replay” a sequence of events to recreate the state of a SQL database at a given time
- ❖ This model presents a lot of possibilities, but it can be challenging to get right
  - particularly when events require participation from external systems

# Event-driven advantages (over REST)

---

## ❖ Asynchronous

- Allows resources to move to the next task once a unit of work is complete
- Events are queued or buffered which prevents consumers from putting back pressure on producers or blocking them

## ❖ Loose Coupling

- Services operate independently, without knowledge of other services, including their implementation details and transport protocol
- Services under an event model can be updated, tested, and deployed independently and more easily

<https://dzone.com/articles/best-practices-for-event-driven-microservice-archi>

# Event-driven advantages (over REST)

---

## ❖ Easy Scaling

- Since the services are decoupled and typically perform only one task, tracking bottlenecks and scaling a service is easier.

## ❖ Recovery Support

- Can recover lost work by “replaying” events from the past.



<https://dzone.com/articles/best-practices-for-event-driven-microservice-archi>

# Event-driven disadvantages

---

- ❖ They are easy to **over-engineer** by separating concerns that might be simpler when closely coupled
  - they can require significant upfront investment, and often result in additional complexity, service contracts or schemas, polyglot build systems, and dependency graphs
- ❖ **Complex data and transaction** management
  - Typically, do not support ACID transactions
  - Systems must carefully handle inconsistent data between services, incompatible versions, duplicate events
- ❖ Even with these drawbacks, ...
  - An event-driven architecture is usually the better choice for enterprise-level microservice systems
  - The pros—scalable, loosely coupled, dev-ops friendly design—outweigh the cons

# Event-driven Anti-Patterns

---

## ❖ Depending on Guaranteed Order and Delivery

- Events are asynchronous
  - Including assumptions of order or duplicates will add complexity and will negate the key benefits of the event-based architecture.

## ❖ Premature Optimization

- Most products start off small and grow over time
  - Consider a simple architecture but include the necessary separation of concerns so that you can swap it out as your needs grow.

## ❖ Expecting Event-Driven to Fix Everything

- Event-driven architecture to fix all the problems
  - It can't fix core problems such as a lack of automated testing, poor team communication, or outdated dev-ops practices.




# Messaging systems

# Messaging systems – models


## ❖ Point to Point – Message queue

- Messages are sent to a queue to pre-defined receivers
- One-to-one relationship between sender and consumer
- Each message is consumed only once



## ❖ Publish-Subscribe

- Each message is published to a topic, and every application that subscribes to that topic gets a copy of all messages published to it.
- Message producers are also known as publishers and consumers are known as subscribers.




# Messaging systems

## ❖ Managing the messages' flow

- Messages are “put into” a source queue
- They are then “taken from” a destination queue
- How/Who/What moves a message from a source queue to a destination queue?

## ❖ Queue Manager / **Message Broker**

- Function as message-queuing “relay” that interact with distributed applications



# Message broker

---

- ❖ Message Broker is built to extend MQ
  - it can understand the content of each message that it moves through the Broker.
- ❖ Message Broker can do the following:
  - divide the publisher and consumer
  - store and route the messages between services
  - converts between different transport protocols
  - identifies and distributes business events from disparate sources


# Message broker

---

- ❖ When is a message broker needed?
  - If we want to **control data feeds**, e.g., the number of registrations in a system
  - When the task is to **put data to several applications** and avoid direct usage of their API
  - When there is a need to **complete processes in a defined order** like a transactional system
- ❖ There are many messaging tools...
  - E.g. Apache ActiveMQ, RabbitMQ
- ❖ ... and protocols
  - E.g. AMQP (Advanced Message Queuing Protocol), MQTT (MQ Telemetry Transport)

# RabbitMQ

- ❖ RabbitMQ is one such open-source message broker software that implements AMQP.

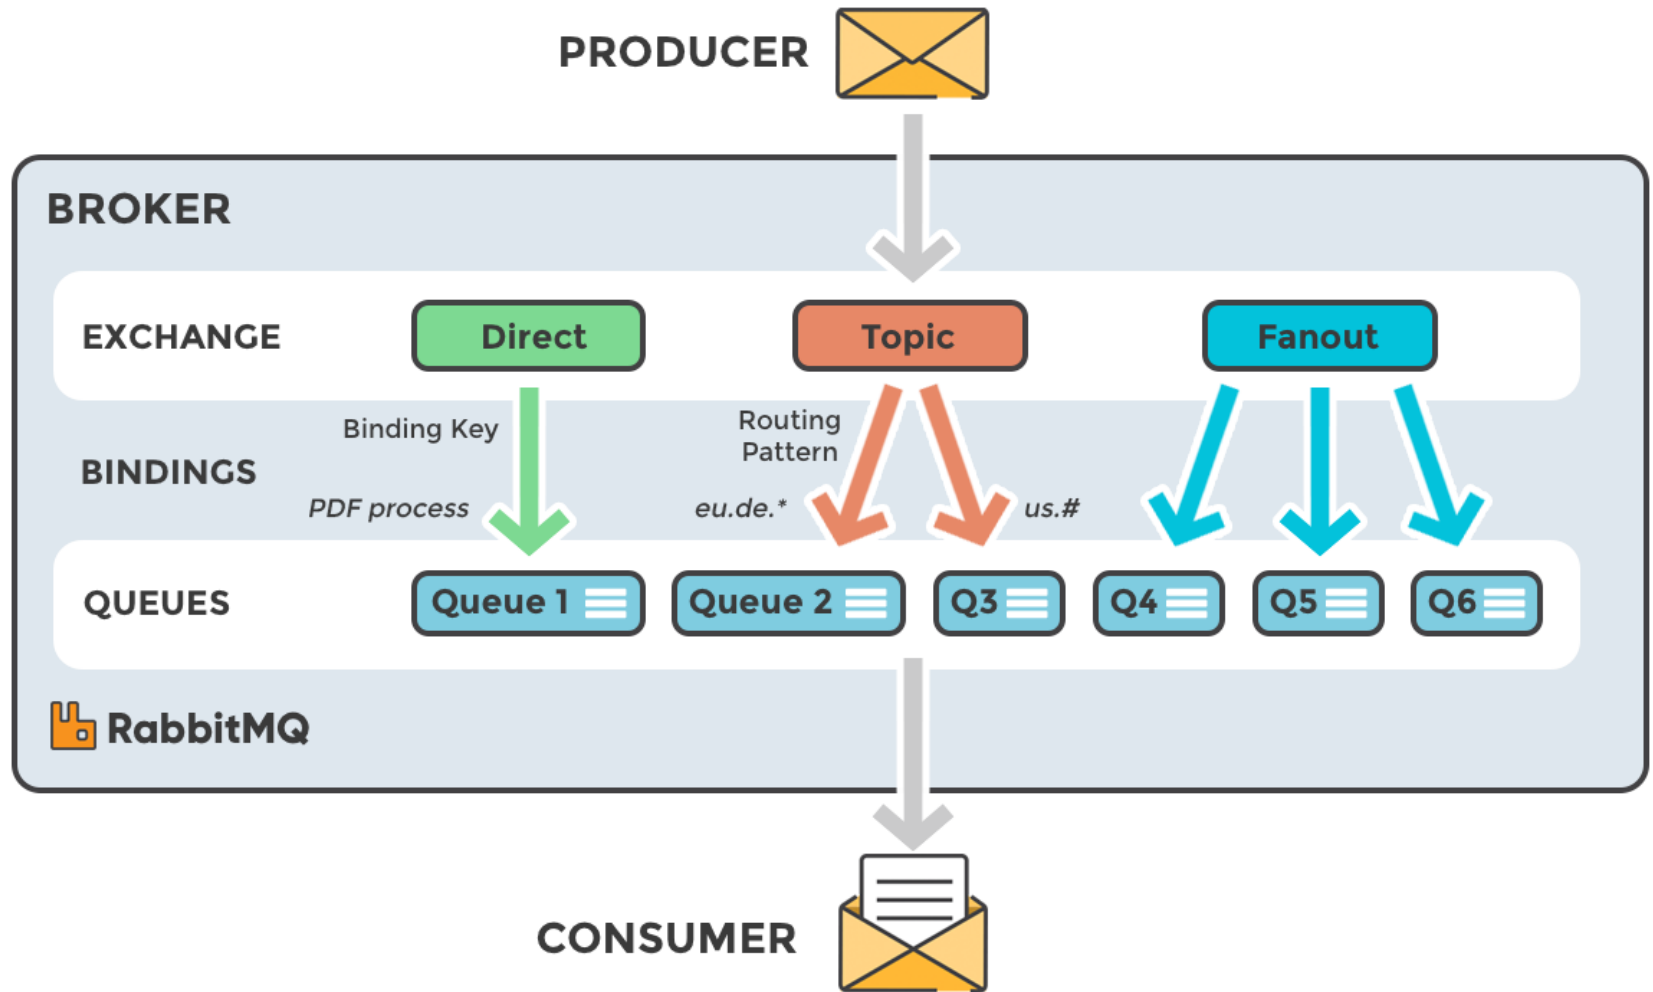


# RabbitMQ – Main concepts

---

- ❖ **Messages** contains attributes (like headers in a request) and a payload (the message content).
- ❖ Messages are published to an entity, **exchange**, which distribute the messages to **queues** (or Topics).
- ❖ The rules for delivering the messages to the right queues are defined through
  - **bindings** (links between exchanges and queues), and
  - **routing keys** (a specific message attribute used for routing).
- ❖ Messages stored in queues are
  - delivered continually to subscribers, or
  - fetched by consumers on demand.

# RabbitMQ – Main concepts





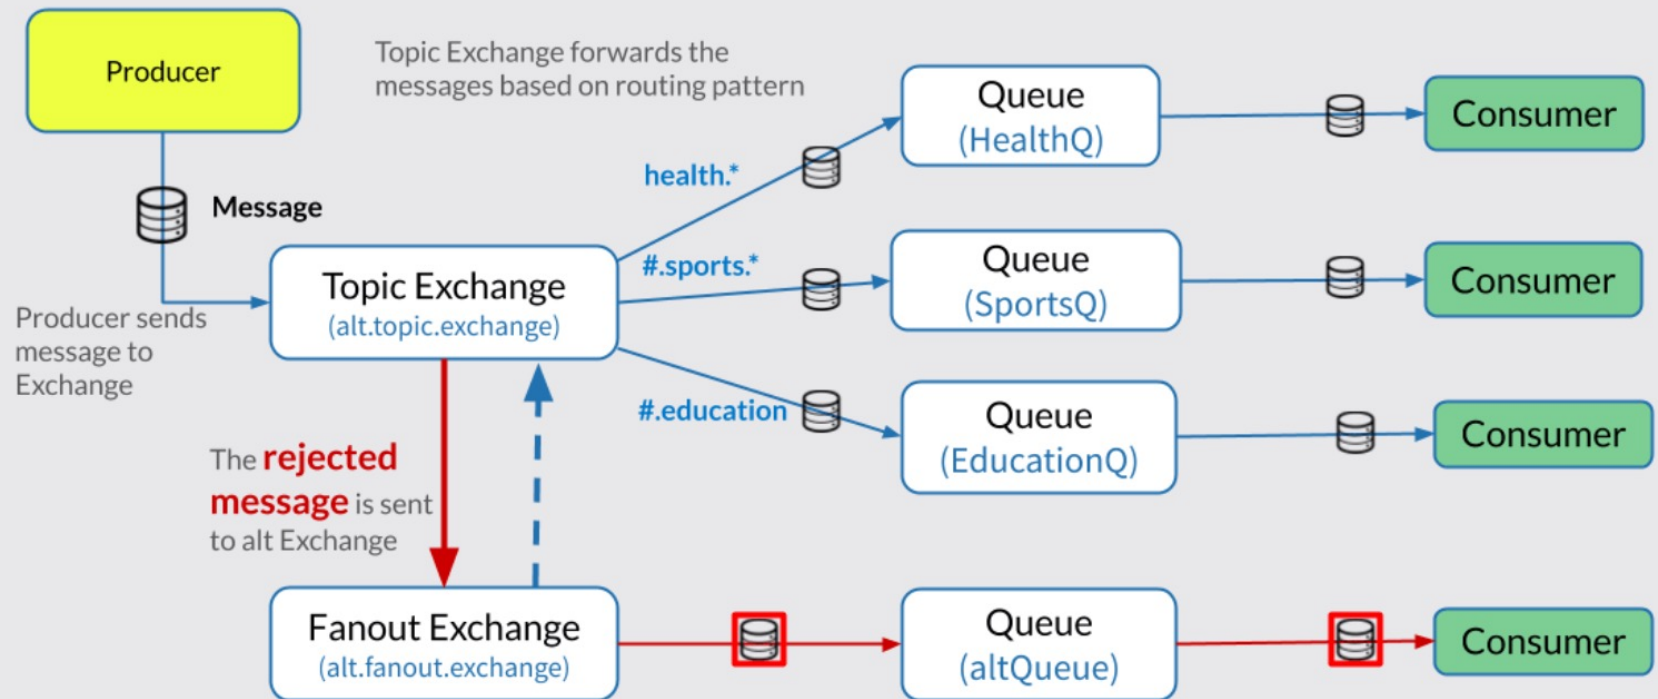
# RabbitMQ – Exchange types (AMQP)

---

- ❖ **Direct Exchange** - It routes messages to a queue by matching **routing key** equal to **binding key**.
- ❖ **Topic Exchange** – It routes messages to multiple queues by a partial matching of a routing key. It uses patterns to match the routing and binding key.
- ❖ **Fanout Exchange** - It ignores the routing key and sends message to all the available queues.
- ❖ **Headers Exchange** – It uses message header instead of routing key.
- ❖ **Default (Nameless) Exchange** - It routes the message to queue name that exactly matches with the routing key.

# RabbitMQ – Exchange types

## Alternate Exchange in RabbitMQ



*jstoligadata.com*

# Spring Boot – Messaging with RabbitMQ

---

- ❖ Set up the RabbitMQ Broker

- ❖ Spring Initializr

- ❖ E.g. RabbitMQ dependency

```
<!-- https://mvnrepository.com/artifact/org.springframework.boot/spring-  
boot-starter-amqp -->  
<dependency>  
<groupId>org.springframework.boot</groupId>  
<artifactId>spring-boot-starter-amqp</artifactId>  
<version>2.3.2.RELEASE</version>  
</dependency>
```

- ❖ <https://spring.io/guides/gs/messaging-rabbitmq/>

# Example excerpt ...

```
@SpringBootApplication
public class MessagingRabbitmqApplication {
    static final String topicExchangeName = "spring-boot-exchange";
    static final String queueName = "spring-boot";

    @Bean
    Queue queue() {
        return new Queue(queueName, false);
    }

    @Bean
    TopicExchange exchange() {
        return new TopicExchange(topicExchangeName);
    }

    @Bean
    Binding binding(Queue queue, TopicExchange exchange) {
        return BindingBuilder.bind(queue).to(exchange).with("foo.bar.#");
    }
}
```

# Hello world – Sender

```
public class Sender {  
  
    @Autowired  
    private RabbitTemplate template;  
  
    @Autowired  
    private Queue queue;  
  
    @Scheduled(fixedDelay = 1000, initialDelay = 500)  
    public void send() {  
        String message = "Hello World!";  
        this.template.convertAndSend(queue.getName(), message);  
        System.out.println(" [x] Sent '" + message + "'");  
    }  
}
```

# Hello world – Receiver

```
@RabbitListener(queues = "hello")
public class Receiver {

    @RabbitHandler
    public void receive(String in) {
        System.out.println(" [x] Received '" + in + "'");
    }
}
```

## ❖ Tutorial:

- <https://www.rabbitmq.com/tutorials/tutorial-one-spring-amqp.html>

# Event streaming systems

## Apache Kafka

# Stored records vs Messaging

## ETL/Data Integration

—

Batch

Expensive

Time Consuming

+

High Throughput

Durable

Persistent

Maintains Order

## Messaging

+

Fast (Low Latency)

—

Difficult to Scale

No Persistence

Data Loss

No Replay

<https://pt.slideshare.net/ConfluentInc/what-is-apache-kafka-and-what-is-an-event-streaming-platform>



# The Event Streaming Paradigm



# The Event Streaming Paradigm

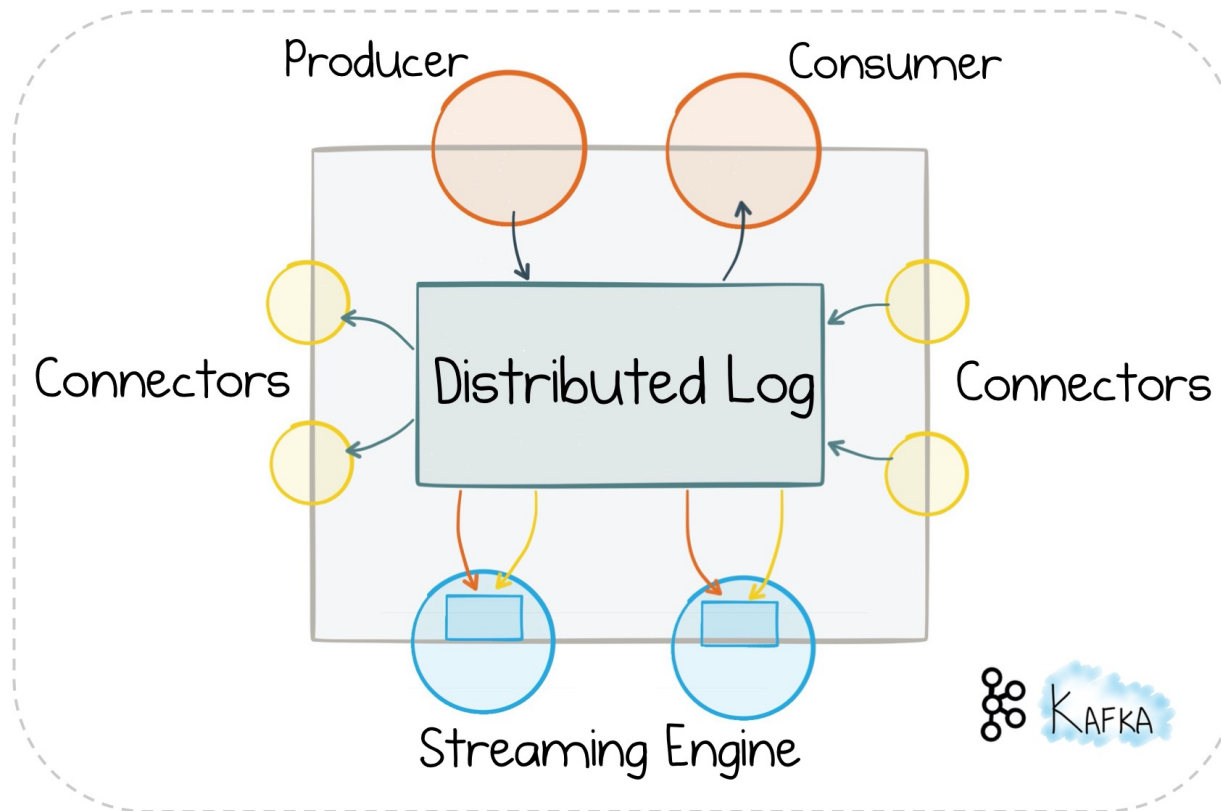
---

***To rethink data as not stored records or transient messages, but instead as a continually updating stream of events***

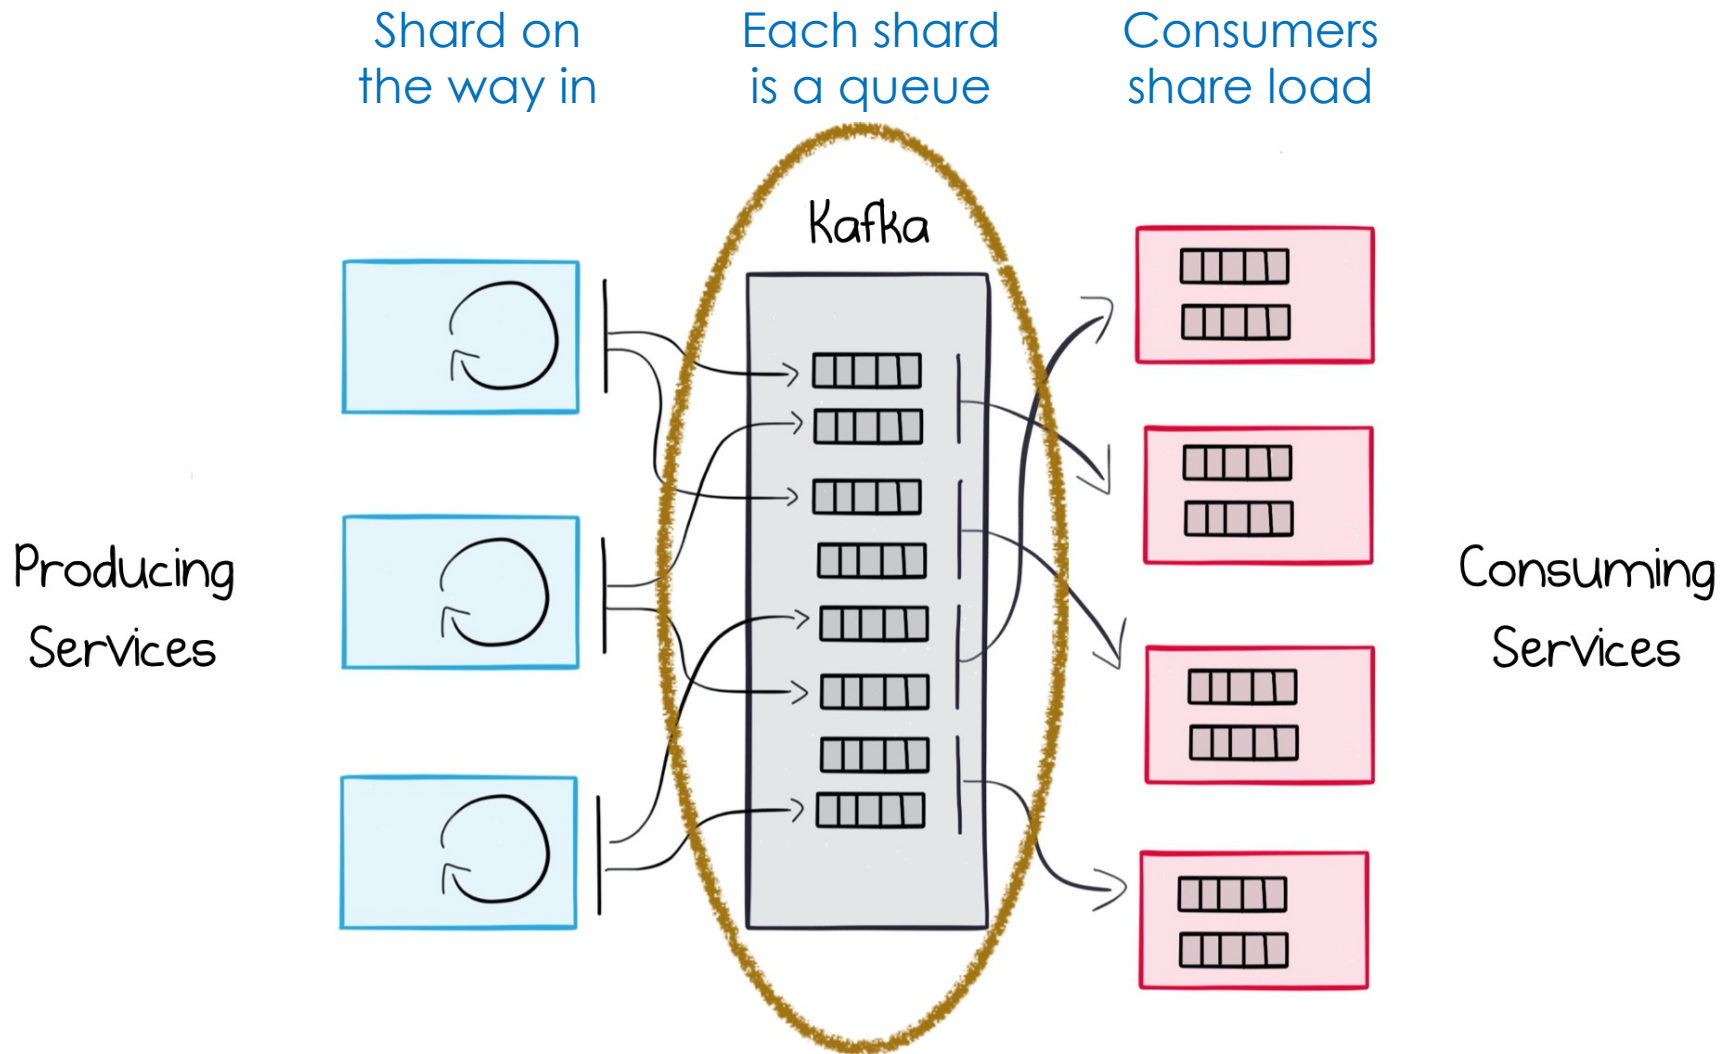
<https://pt.slideshare.net/ConfluentInc/what-is-apache-kafka-and-what-is-an-event-streaming-platform>

# Apache Kafka

- ❖ Apache Kafka is made of distributed, immutable, append-only commit logs





# The Distributed Log



# The Distributed Log

## ❖ The Service Backbone

- Scalable/Load Balanced, Fault Tolerant, Concurrent, Strongly Ordered, Stateful



## ❖ A place to keep data-on-the-outside

- When sending data across services, it is outside the normal "trust" boundary.

# Events and topics

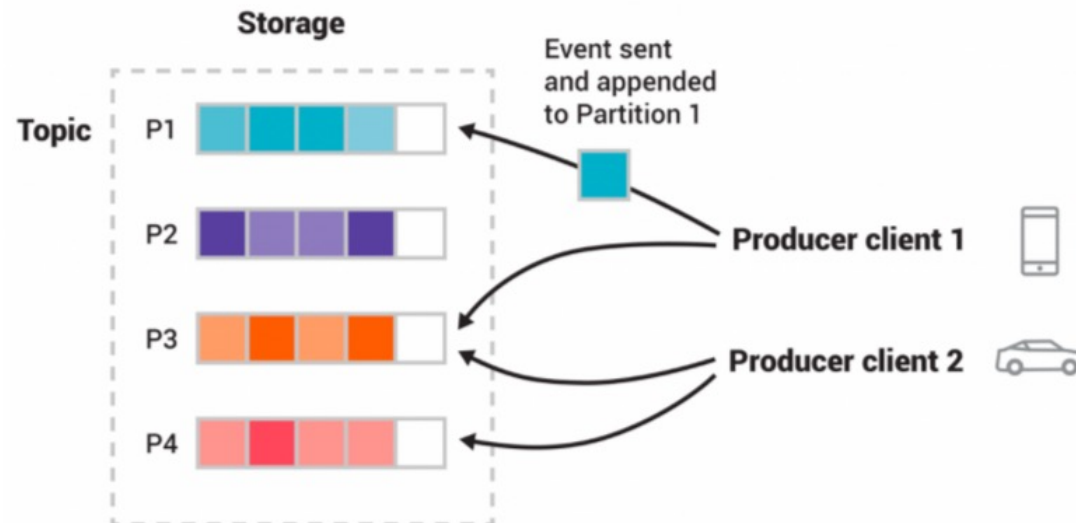
- ❖ Each **event** has a key, value, timestamp, and optional metadata headers.

Event key: "Alice"

Event value: "Made a payment of \$200 to Bob"

Event timestamp: "Jun. 25, 2020 at 2:06 p.m."

- ❖ Events are organized and durably stored in **topics**
- ❖ Topics are **partitioned**
  - a topic is spread over "buckets" located on different Kafka brokers.



# Kafka APIs

---

- ❖ **Admin API** to manage and inspect topics, brokers, and other Kafka objects.
- ❖ **Producer API** to publish (write) a stream of events to one or more Kafka topics.  
`KafkaProducer<K,V>` (implements `Producer<K,V>`)
- ❖ **Consumer API** to subscribe to (read) one or more topics and to process the stream of events produced to them.  
`KafkaConsumer<K, V>` (implements `Consumer<K,V>`)
- ❖ The **Kafka Streams API** to implement stream processing applications and microservices.  
`KStream<K, V>`
  - To process event streams, including transformations, stateful operations like aggregations and joins, windowing, processing based on event-time, and more.
- ❖ The **Kafka Connect API** to build and run import/export connectors to external systems and applications.
  - Kafka community already provides hundreds of ready-to-use connectors.

# Event-driven patterns

---

- ❖ Previously...
  - Event notification, Event-carried state transfer, Event-sourcing
- ❖ In **Event Sourcing**, events are a core element – the source of truth.
  - Being stored, immutably, in the order they were created in, the event log expresses exactly what the system did.
- ❖ **Command Query Response Segregation (CQRS)** is a natural progression from this.
  - Decoupling writing and reading operations.
  - As a simple example, we might write events to Kafka (write model), read them back, and then push them into a database (read model).




# Example

- ❖ Say we have an email service that listens to an event stream of orders and then sends confirmation emails to users once they complete a purchase.
  - This requires information about both the order as well as the associated payment. Such an email service might be created in several different ways



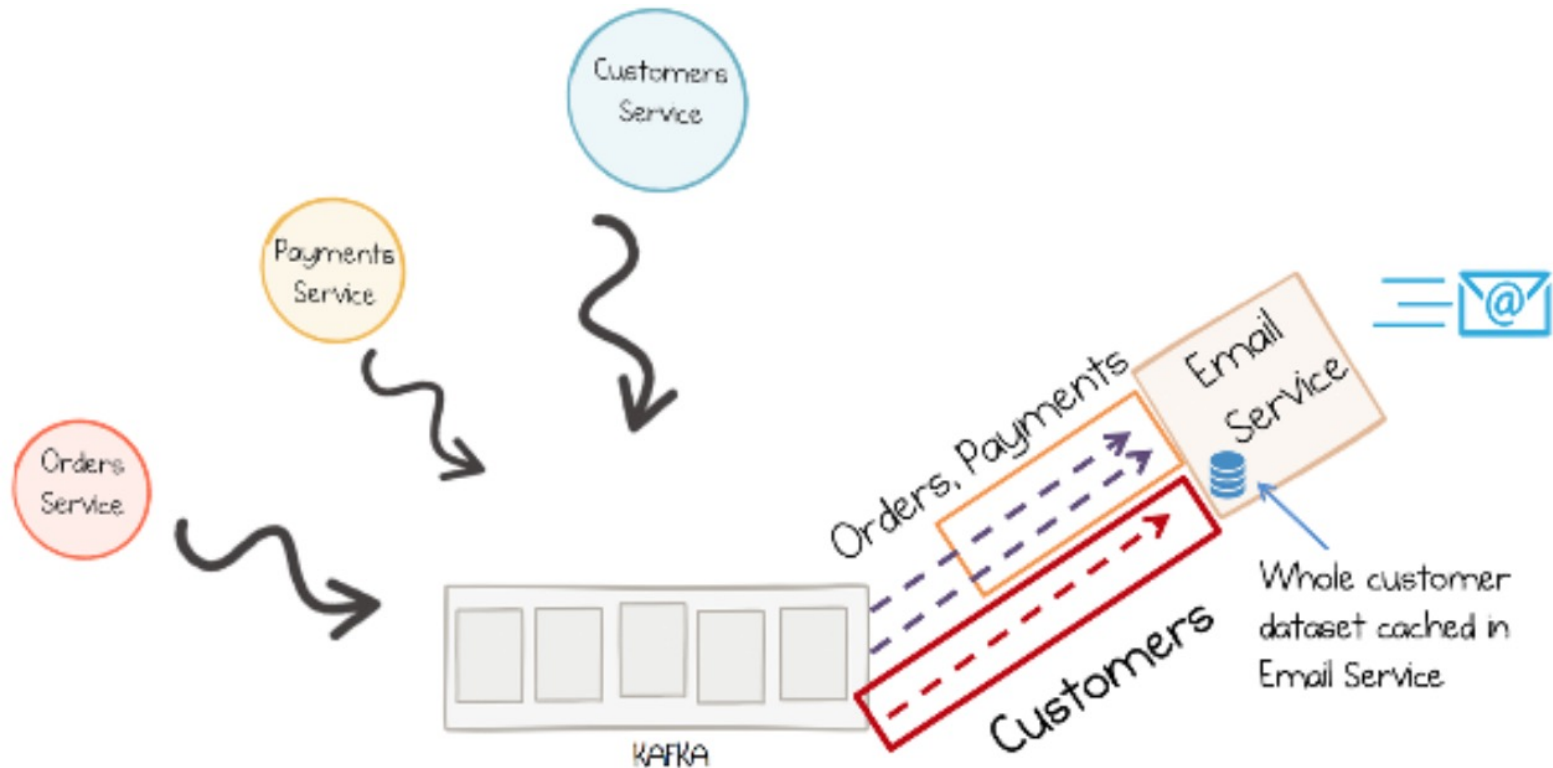
# Stateless Streaming approach

- ❖ Example: A **stateless streaming** service that looks up reference data in another service at runtime



# Stateful Streaming approach

- ❖ Example: A **stateful streaming** service that replicates the Customers topic into a local table, held inside the Kafka Streams API.




# Stateful Streaming approach

---

- ❖ Being stateful comes with some challenges:
  - when a new node starts, it must load all stateful components (i.e., state stores)
- ❖ **Kafka Streams**, for instance, provides three mechanisms to simplify stateful:
  - It uses a technique called **standby replicas**, which ensure that for every table or state store on one node, there is a replica kept up to date on another.
  - **Disk checkpoints** are created periodically so that, should a node fail and restart, it can load its previous checkpoint.
  - Finally, **compacted topics** are used to keep the dataset as small as possible. This acts to reduce the load time for a complete rebuild should one be necessary.


# A Stateful Stream Processing



Database

Finite source

VS



Stateful Stream Processor

Infinite & Finite source

# Messaging in Spring

---

- ❖ **Spring Cloud Stream** helps fully abstract code from the underlying messaging engine
- ❖ It supports a variety of binder implementations such as:
  - RabbitMQ
  - Apache Kafka
  - Kafka Streams
  - Amazon Kinesis
  - Google PubSub (partner maintained)
  - Solace PubSub+ (partner maintained)
  - Azure Event Hubs (partner maintained)
  - Apache RocketMQ (partner maintained)

# Messaging in Spring

---

## ❖ Spring AMQP

- It applies core Spring concepts to the development of AMQP-based messaging solutions.
- The project consists of two parts; spring-amqp is the base abstraction, and spring-rabbit is the RabbitMQ implementation.

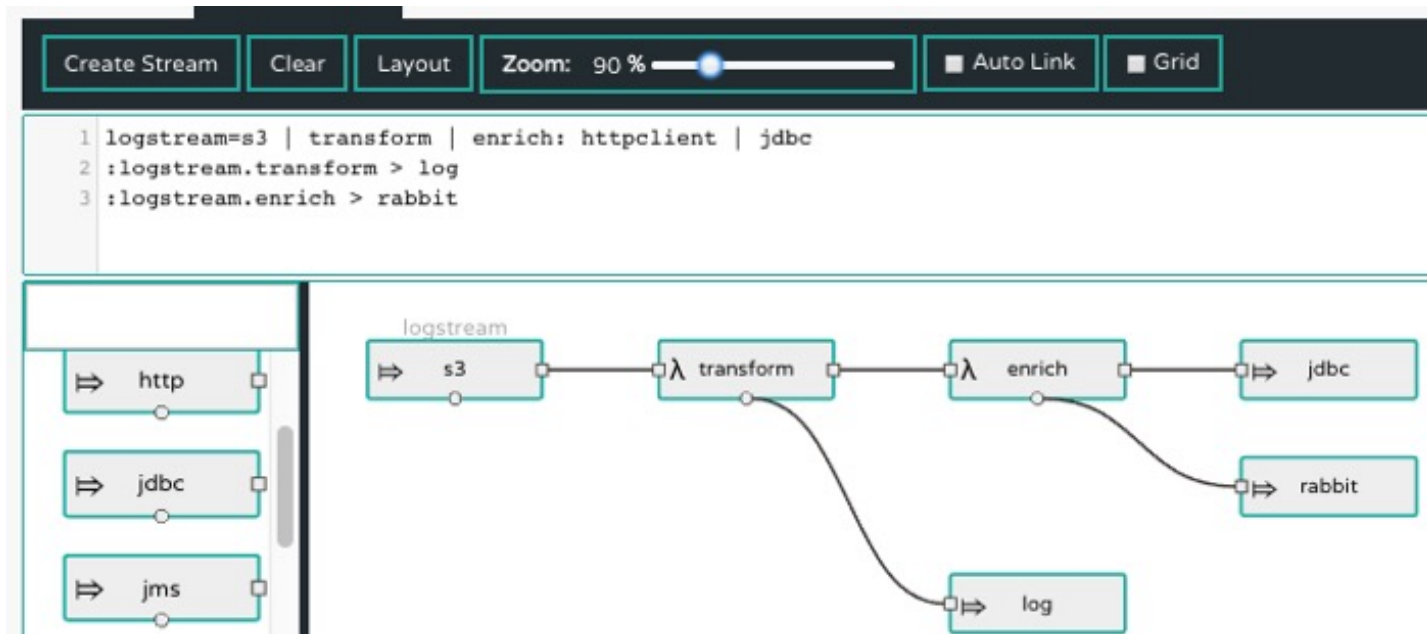
## ❖ Spring for Apache Kafka

- Simple snippet

```
@Configuration(proxyBeanMethods = false)
@EnableKafkaStreams
public static class KafkaStreamsExampleConfiguration {
    @Bean
    public KStream<Integer, String> kStream(StreamsBuilder streamsBuilder) {
        KStream<Integer, String> stream = streamsBuilder.stream("ks1In");
        stream.map((k, v) -> new KeyValue<>(k, v.toUpperCase())).to("ks1Out",
            Produced.with(Serdes.Integer(), new JsonSerde<>()));
        return stream;
    }
}
```

# Messaging in Spring

- ❖ **Spring Cloud Data Flow** allows to create and orchestrate data pipelines, e.g., data ingest, real-time analytics, and data import/export.






# Summary

---

- ❖ Event-driven architecture is gaining in popularity, and with good reason.
  - From a technical perspective, it provides an effective method of wiring up microservices.
  - The interest in serverless functions - such as AWS Lambda, Azure Functions, or Knative - is growing, and these are inherently event-based.
  - Moreover, when coupled with modern streaming data tools like Apache Kafka, event-driven architectures become more versatile, resilient, and reliable than with earlier messaging methods.
- ❖ But perhaps the most important “feature” of the event-driven pattern is that it models how businesses operate in the real world.

# Resources & Credits



- ❖ Designing Event-Driven Systems  
Ben Stopford, O'Reilly
- ❖ The Optimal RabbitMQ Guide, Lovisa Johansson, CloudAMQP
- ❖ Spring messaging projects
  - <https://spring.io/projects/>
- ❖ Kafka, RabbitMQ, etc..
  - <https://kafka.apache.org>
  - <https://www.rabbitmq.com>